

UNDERSTANDING OPTIMIZATION *in* REINFORCEMENT LEARNING



*An Empirical Study of
Algorithms and their Hyperparameters*

Jan Ole von Hartz

May 2019

*Submitted in partial fulfillment of the requirements
for the degree of Bachelor of Science*

to the

*Department of Computer Science
Technical Faculty
University of Freiburg*

This thesis was typeset using the incredible template created and released into the public domain by Eivind Uggedal, found at <https://github.com/uggedal/thesis>.

Hereafter follow the original remarks to his thesis.

This thesis was typeset using the L^AT_EX typesetting system originally developed by Leslie Lamport, based on T_EX created by Donald Knuth.

The body text is set 12/14.5pt on a 26pc measure with Minion Pro designed by Robert Slimbach. This neohumanistic font was first issued by Adobe Systems in 1989 and have since been revised. Other fonts include Sans and Typewriter from Donald Knuth's Computer Modern family.

Typographical decisions were based on the recommendations given in *The Elements of Typographic Style* by Bringhurst (2004).

The use of sidenotes instead of footnotes and figures spanning both the textblock and fore-edge margin was inspired by *Beautiful Evidence* by Tufte (2006).

The guidelines found in *The Visual Display of Quantitative Information* by Tufte (2001) were followed when creating diagrams and tables. Colors used in diagrams and figures were inspired by the *Summer Fields* color scheme found at <http://www.colourlovers.com/palette/399372>



DECLARATION

I hereby declare, that I am the sole author and composer of my thesis and that no other sources or learning aids, other than those listed, have been used. Furthermore, I declare that I have acknowledged the work of others by providing detailed references of said work. I also hereby declare that my thesis has not been prepared for another examination or assignment, either in its entirety or excerpts thereof.

Place, date

Signature

ABSTRACT

In machine learning, praxis is by far more delicate than theory: not just different algorithms and random seeds, but also the usage of different optimizers and settings of their hyperparameters yield dramatically different results, to the point of convergence versus non-convergence. Deep reinforcement learning is especially difficult due to the strongly moving loss landscape, with agents not learning the desired behavior because of getting stuck in local optima. While there exists some rough ideas, the exact effects of different optimizers (such as Adam) and their hyperparameters on the returns of the agent, as well as the stability in training and evaluation, are not yet well understood. Using the well known DQN algorithm for discrete and DDPG algorithm for continuous environments, we investigate the influence of different optimizers and settings of their hyperparameters on the agents performance, as well as compare the optimizers' sensitivity to hyperparameters and their stability in training. Using BOBH, a hyperparameter optimization method, we further try to make a step towards hyperparameter agnostic deep reinforcement learning, to try to avoid the pitfalls of hyperparameter settings.

CONTENTS

Abstract i

Contents iii

List of Figures v

List of Tables vi

Preface vii

1 Introduction 1

Background

- 2 Reinforcement Learning 5
 - 2.1 The Mathematical Model 5
 - 2.2 Application of the Model 6
- 3 Algorithms and Methods 9
 - 3.1 Numerical Optimization 9
 - 3.2 Agents 9
 - 3.3 Hyperparameter Optimization 9
- 4 Tools and Libraries 11
 - 4.1 Numerical Optimization 11
 - 4.2 Agents 11
 - 4.3 Environments 11
 - 4.4 Hyperparameter Optimization 11
 - 4.5 Visualization and Evaluation 11

Performance Estimation

- 5 Problem 15
- 6 Experimental Setup 17
- 7 Results 19

8	Discussion	21
---	------------	----

Optimizer Analysis

9	Problem	25
10	Experimental Setup	27
11	Results	29
12	Discussion	31

Summary

13	Conclusion	35
	Bibliography	37

Appendices

LIST OF FIGURES

LIST OF TABLES

PREFACE

none

JAN OLE VON HARTZ
Freiburg, Germany
May 2019

INTRODUCTION

According to Sutton and Barto (2018) the field of *reinforcement learning* studies learning “how to map situations to actions—so as to maximize a numerical reward signal”, and refers - besides the field of study - to both to the problem class itself, as well as a class of solution methods. The usually used formalism is that of an *agent* interacting with an *environment* via a fixed set of *actions* and only observing 1. a numerical reward signal and 2. the state of the environment or an observation which describes part of it. In recent years, *reinforcement learning* has not only seen a surge of research, but has also produced a number of remarkable results, such as programs learning to control a helicopter (Ng et al., 2006), locomotion tasks (Heess et al., 2017), in-hand manipulation of objects (OpenAI et al., 2018) and playing games such backgammon (Tesauro, 1995) Atari video games (Mnih et al., 2013) and Go (Silver et al., 2016, 2017). This is in part due to the increasing application of so called *deep reinforcement learning* techniques, the use of *deep neural networks* in reinforcement learning. While allowing for great scientific progress, the use of deep neural networks also introduces many new and opaque problems, some of which we try to tackle in thesis. For one, solving machine learning problems usually involves the selection of a as well suited as possible function from a class of functions by minimizing some loss function defined over the output of said function for a set of data points. Hence, the training of a deep neural network has in its core a *numerical optimization problem*: finding a suited parametrization of the network to achieve the best possible performance (the lowest possible loss) on unseen data by using a given data set as an approximation of the general data distribution and optimizing the networks performance on it. This is usually done via *gradient-based methods*, such *gradient descent* and variations of it. Reinforcement learning however introduces additional problems; unlike in *supervised learning*, the training data is not given, but must be produced by agent by engaging with the environment in a trial-and-error fashion. Since the deep neural networks are used by the agent to e.g. approximate a function that estimates the reward of a certain action in a certain state and these estimations change drastically over time, the numerical optimization of these networks is very unstable. A systematic problem of gradient-based methods is that unlike analytical methods they are only suited to find local optima, not global optima. In reinforcement learning the loss landscape is usually very rough and provides plenty of local optima to get stuck in. Furthermore, it is not



well understood, how different gradient-based methods compare in reinforcement learning, especially since these are also very sensitive to the settings of their hyperparameters. A common call thus is for the development of hyperparameter agnostic algorithms - algorithms that adapt their hyperparameters themselves during runtime (Henderson et al., 2018). For this thesis we

1. compared different gradient-based optimizers in terms of
 - a) the final performance of their produced network configuration
 - b) the stability of these optima
 - c) their sensitivity towards their hyperparameters
 - d) their stability in training

to try to get an understanding of their characteristics in a in deep reinforcement learning setting.

2. Using a hyperparameter optimization method, automatically find well performing hyperparameter settings for the optimizers to make a step towards hyperparameter agnosticism.

The automatic hyperparameter optimization introduced further challenges, since it requires for a cost-effective estimation of the performance of a given configuration - something that is not given in reinforcement learning, where the training of an agent is usually very costly.

Next, we will introduce the necessary background in a logical order; starting with the mathematical foundation, followed by the used algorithms and finally their implementation.

PART I

BACKGROUND

REINFORCEMENT LEARNING

As introduced, problems in reinforcement learning are usually framed as an *agent* interacting with an *environment* via a fixed set of *actions* and only observing 1. a numerical reward signal and 2. the state of the environment or an observation which describes part of it. In this chapter introduces the mathematical formalism and the theoretical ideas behind its practical application.

2

2.1 THE MATHEMATICAL MODEL

In large parts of the current research, as well as this thesis, instances of this problem class are formalized using *Markov Decision Processes* (Howard, 1960). We thus now define this and related notions.

2.1.1 Markov Decision Processes and Policies

Definition 1 (Markov Process) A Markov Process is a 2-tuple $(\mathcal{S}, \mathcal{P})$, where \mathcal{S} is a finite set of states and \mathcal{P} is a state transition function $\mathcal{P}(s, s') = \mathbf{P}[s_t = s' \mid s_{t-1} = s]$, that fulfills the Markov Property, i.e. for all states and at all times $\mathbf{P}[s_t \mid s_{t-1}, s_{t-2}, \dots, s_0] = \mathbf{P}[s_t \mid s_{t-1}]$.

Definition 2 (Markov Decision Process) A Markov Decision Process is a 5-tuple $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$, where $(\mathcal{S}, \mathcal{P})$ is a Markov Process with transition function $\mathcal{P}_a(s, s') = \mathbf{P}[s_t = s' \mid s_{t-1} = s, a_{t-1} = a]$ (where $a \in \mathcal{A}$), \mathcal{A} is a finite set of actions, $\gamma \in [0, 1]$ is a discount factor and \mathcal{R} is the expected immediate reward $\mathcal{R}_a(s, s') = \mathbf{E}[R_t \mid s_{t-1} = s, s_t = s', a_{t-1} = a]$, with R_t being the reward received at point t , which can be a function mapping s_{t-1} , s_{t-1} and s_t or s_{t-1} and s_t and a_{t-1} to some real number.

Definition 3 (Policy) A policy π is a mapping from a state space \mathcal{S} to a set of probability distributions over an action space \mathcal{A} given some state $s \in \mathcal{S}$: $s \mapsto p(\mathcal{A} = a \mid s)$.

Reinforcement learning problems are then often modeled as interpreting a *Markov Decision Process* M as an environment in which an Agent A can take Actions $a \in \mathcal{A}_M$ and observe the ensuing reward $r \in \mathbf{R}$ and state $s \in \mathcal{S}_M$. Its goal then usually is to find an optimal policy π^* to maximize some cumulative function, usually the total discounted reward $R = \sum_{t=0}^{T-1} \gamma^t R_{a_t}(s_t, s_{t+1})$, where T is called the *horizon* of the problem.

2.1.2 Value Functions and the Bellmann Equation

Notably, a policy implies a transition distribution over the state space of a Markov Decision Process M . This property allows one to easily evaluate a given policy π on M by reducing the problem of evaluation to the evaluation of a Markov Process, like the following definition illustrates.

Definition 4 (State-Value Function) *The state-value function $V^\pi(s)$ of a state $s \in \mathcal{S}$ given a policy π is the expected return gained when starting in s and following π henceforth: $V^\pi(s) = \mathbb{E}_\pi[\sum_{t=0}^{T-1} \gamma^t R_{a_t}(s_t, s_{t+1}) \mid s_0 = s]$.*

Definition 5 (State-Action-Value Function) *The state-action-value function $Q^\pi(s, a)$ of a state $s \in \mathcal{S}$ and an action $a \in \mathcal{A}$ given a policy π is the expected return gained when starting in s , taking a and following π thereafter: $Q^\pi(s, a) = \mathbb{E}_\pi[\sum_{t=0}^{T-1} \gamma^t R_{a_t}(s_t, s_{t+1}) \mid s_0 = s, a_0 = a]$.*

The *state-action-value function* (often called *Q-function*) fulfills the following recursive property (a similar property holds for the *action-value function*):

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_a(s, s') \sum_{a' \in \mathcal{A}} \pi(a' \mid s') \cdot Q^\pi(s', a') \quad (\text{Bellmann Equation})$$

This famous property called *Bellmann Equation* allows to, given the reward function R , iteratively compute the Q -function (or V -function) (Bellman, 1957). Remember, that the goal for the agent is to learn the optimal policy π^* , which corresponds to an optimal *state-action-value function* Q^* . The optimal policy π^* is the one that maximizes the corresponding Q -function Q^* . As the *Bellmann Equation* illustrates, maximizing Q under policy π is achieved by greedily selecting the action a in state s that maximizes $R(s, a) + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_a(s, s') \sum_{a' \in \mathcal{A}} \pi(a' \mid s') \cdot Q^\pi(s', a')$, yielding an improved policy π' : $\pi'(s, a) = \begin{cases} 1 & \text{if } a = \arg \max_{\hat{a} \in \mathcal{A}} Q(s, \hat{a}) \\ 0 & \text{else} \end{cases}$.

Hence, just as Q^* can iteratively be computed, π^* can be iteratively computed as well, e.g. via *dynamic programming* (Howard, 1960).

2.2 APPLICATION OF THE MODEL

These lessons learned in the exact solution of the calculation of the optimal policy via dynamic programming can now also be applied to approximately solving this problem, as usually done in current reinforcement learning research. For this, we introduce two approaches that the algorithms that we introduce later make use of: *Value Base Methods* and *Actor-Critic Methods*.

2.2.1 *Value Based and Actor-Critic Methods*

To learn an approximation to the optimal policy π^* in an environment M an agent A can thus estimate Q^* via bootstrapping and iteratively improving its estimates, which is at the heart of so called *Value Based Methods*, such as *Q-learning* (Watkins and Dayan, 1992).

As the selection of the policy often happens in a greedy manner, all possible actions must be considered, which is only feasible for action spaces that are discrete and finite. For other action spaces, agents can for example separately estimate both the V -function and the policy (instead of producing the current policy via greedy selection from the current estimation of Q), which is at the heart of so called *actor-critic methods*. Another issue with continuous action spaces is that the runtime of the iterative update of the Q -function is at least linearly dependent on the size of the action space (see Bellmann Equation) . This can be solved by approximating the updates of the Q -function via *Monte Carlo sampling* from the action space, which we will not discuss. Another common approach is to artificially discretize a continuous action space.

2.2.2 *Practicality of the Markov Property*

As noted by Arulkumaran et al. (2017), the underlying assumption of the *Markov Property* does often not hold in practical applications, since it requires full observability of the states. While algorithms that use *Partly Observable Markov Decision Processes* exist, full observability is often simply assumed as an approximation to the actual partly observable environment. An example of such an approximation can be found in Mnih et al. (2013); using the video output of an Atari game simulator the authors preprocessed the observations by combining four video frames into a single observation to approximate state features such as velocity of game objects.

We will now introduce the used algorithms and explain how they relate to the introduced formalisms and methods in order of increasing abstraction level.

ALGORITHMS AND METHODS

3.1 NUMERICAL OPTIMIZATION

3.2 AGENTS

3.3 HYPERPARAMETER OPTIMIZATION

3

TOOLS AND LIBRARIES

Ref to henderson about code base influence etc.

4.1 NUMERICAL OPTIMIZATION

4.2 AGENTS

4.3 ENVIRONMENTS

4.4 HYPERPARAMETER OPTIMIZATION

4.5 VISUALIZATION AND EVALUATION



PART II

PERFORMANCE ESTIMATION

PROBLEM

5

EXPERIMENTAL SETUP

6

RESULTS

7

DISCUSSION

8

PART III

OPTIMIZER ANALYSIS

PROBLEM

9

EXPERIMENTAL SETUP

10

RESULTS

11

DISCUSSION

12

PART IV

SUMMARY

CONCLUSION

None

13

BIBLIOGRAPHY

- Arulkumaran, Kai; Deisenroth, Marc Peter; Brundage, Miles; and Bharath, Anil Anthony. 2017. *A brief survey of deep reinforcement learning*. In arXiv preprint arXiv:1708.05866. Cited on p. 7.
- Bellman, Richard. 1957. *A Markovian decision process*. In Journal of Mathematics and Mechanics, vol. 6, no. 5, pp. 679–684. Cited on p. 6.
- Brigham, Robert. October 2004. *The elements of typographic style*. Hartley & Marks Publishers, Point Roberts, WA, USA, 3rd edn. ISBN 0-881-79205-5. Cited on p. b.
- Heess, Nicolas; Sriram, Srinivasan; Lemmon, Jay; Merel, Josh; Wayne, Greg; Tassa, Yuval; Erez, Tom; Wang, Ziyu; Eslami, SM; Riedmiller, Martin; et al. 2017. *Emergence of locomotion behaviours in rich environments*. In arXiv preprint arXiv:1707.02286. Cited on p. 1.
- Henderson, Peter; Islam, Riashat; Bachman, Philip; Pineau, Joelle; Precup, Doina; and Meger, David. 2018. *Deep reinforcement learning that matters*. In Thirty-Second AAAI Conference on Artificial Intelligence. Cited on p. 2.
- Howard, Ronald A. 1960. *Dynamic Programming and Markov Processes*. In . Cited on pp. 5 and 6.
- Mnih, Volodymyr; Kavukcuoglu, Koray; Silver, David; Graves, Alex; Antonoglou, Ioannis; Wierstra, Daan; and Riedmiller, Martin. 2013. *Playing atari with deep reinforcement learning*. In arXiv preprint arXiv:1312.5602. Cited on pp. 1 and 7.
- Ng, Andrew Y; Coates, Adam; Diel, Mark; Ganapathi, Varun; Schulte, Jamie; Tse, Ben; Berger, Eric; and Liang, Eric. 2006. *Autonomous inverted helicopter flight via reinforcement learning*. In Experimental Robotics IX, pp. 363–372. Springer. Cited on p. 1.
- OpenAI; Andrychowicz, Marcin; Baker, Bowen; Chociej, Maciek; Józefowicz, Rafal; McGrew, Bob; Pachocki, Jakub W.; Pachocki, Jakub; Petron, Arthur; Plappert, Matthias; Powell, Glenn; Ray, Alex; Schneider, Jonas; Sidor, Szymon; Tobin, Josh; Welinder, Peter; Weng, Lilian; and Zaremba, Wojciech. 2018. *Learning Dexterous In-Hand Manipulation*. In CoRR, vol. abs/1808.00177. URL <http://arxiv.org/abs/1808.00177>. Cited on p. 1.

- Silver, David; Huang, Aja; Maddison, Chris J; Guez, Arthur; Sifre, Laurent; Van Den Driessche, George; Schrittwieser, Julian; Antonoglou, Ioannis; Panneershelvam, Veda; Lanctot, Marc; et al. 2016. *Mastering the game of Go with deep neural networks and tree search*. In *nature*, vol. 529, no. 7587, p. 484. Cited on p. 1.
- Silver, David; Schrittwieser, Julian; Simonyan, Karen; Antonoglou, Ioannis; Huang, Aja; Guez, Arthur; Hubert, Thomas; Baker, Lucas; Lai, Matthew; Bolton, Adrian; et al. 2017. *Mastering the game of go without human knowledge*. In *Nature*, vol. 550, no. 7676, p. 354. Cited on p. 1.
- Sutton, Richard S and Barto, Andrew G. 2018. *Reinforcement Learning: An Introduction*. The MIT Press, Cambridge, MA, USA, 2nd edn. ISBN 978-0-26203-924-6. Cited on p. 1.
- Tesauro, Gerald. 1995. *Temporal difference learning and TD-Gammon*. In *Communications of the ACM*, vol. 38, no. 3, pp. 58–69. Cited on p. 1.
- Tufte, Edward R. may 2001. *The Visual Display of Quantitative Information*. Graphics Press LLC, Cheshire, CT, USA, 2nd edn. ISBN 0-961-39214-2. Cited on p. b.
- . jul 2006. *Beautiful Evidence*. Graphics Press LLC, Cheshire, CT, USA. ISBN 0-961-39217-7. Cited on p. b.
- Watkins, Christopher JCH and Dayan, Peter. 1992. *Q-learning*. In *Machine learning*, vol. 8, no. 3-4, pp. 279–292. Cited on p. 7.

APPENDICES

none